



Music Generation

Research style project:
An exploration of Melspectrogram.

cMelGAN - conditional generative model based
on MelSpectrograms, a faster model for limited
hardware computing resources

Digital Representation of Music

- Sound is a **wave**

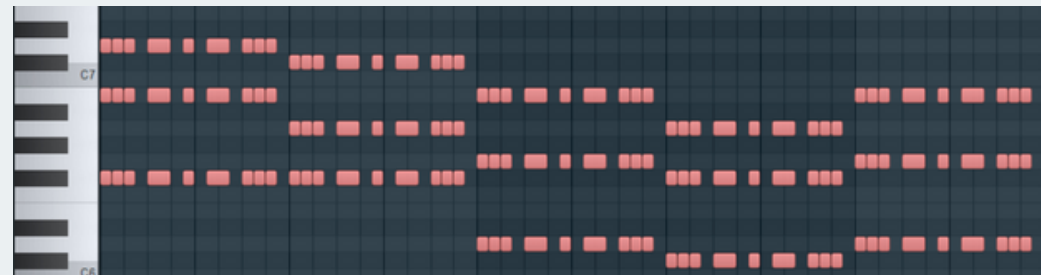
- Sheet music of **notes**

- Present music in **time** and **frequency** domain
- Closest to **human perception** of music


WAV form




MIDI



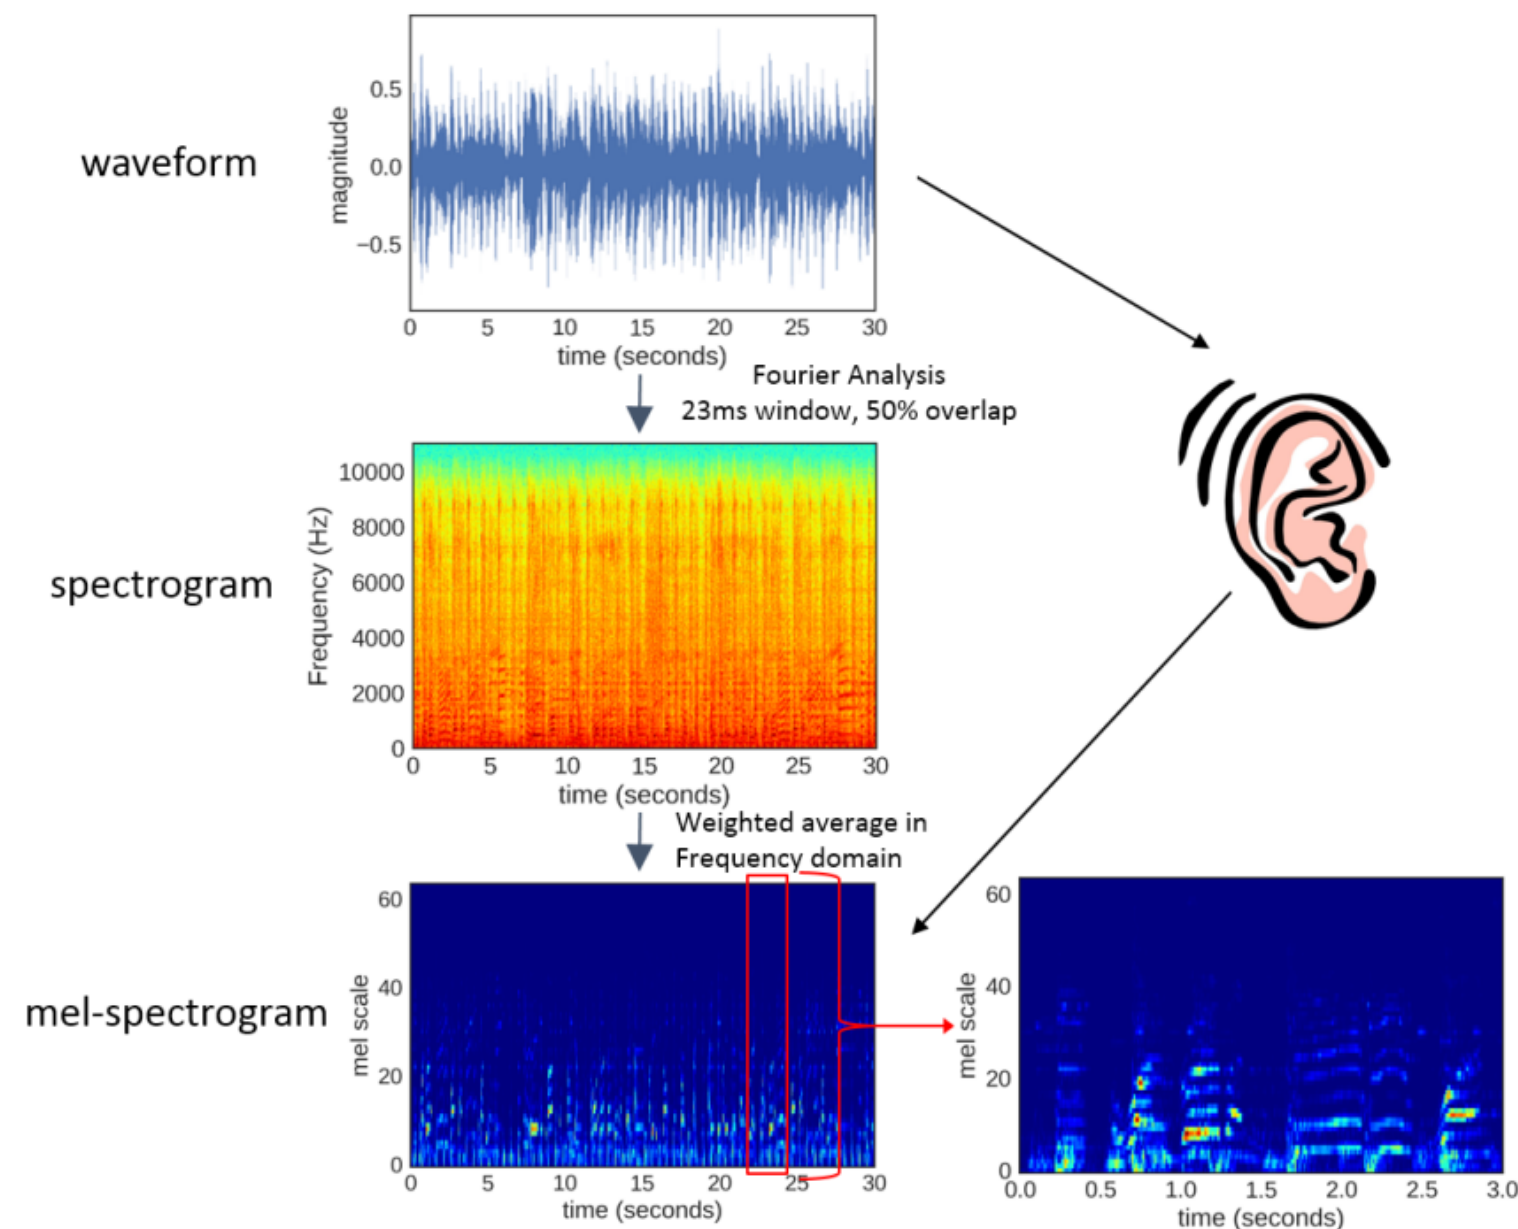
MelSpectrogram



Mel Spectrogram pictures of sounds



- 2-D representation of audios by taking **fourier transform** of the waveform at regular interval.
- Human cochlea contains hairs associated with **one particular frequency** to extract the frequency information given an audio **wave**.
- Would this representation generate **better** music?



Fourier analysis
happening right now...

Problem Statement



PROBLEM
STATEMENT

PASTWORK

DATACOLLECTION

- Can Mel Spectrograms **outperform** traditional MIDI-based formats?
- Add genre **conditioning** because its **rarely** been done in generative models based on MelSpectrograms
- Potentially **closer** to human perception
- Downside is **less structured** format so harder for network to learn good representations
- Maybe similar to how convolutional networks match human perception of images and so do better on image-based tasks

PAST WORK

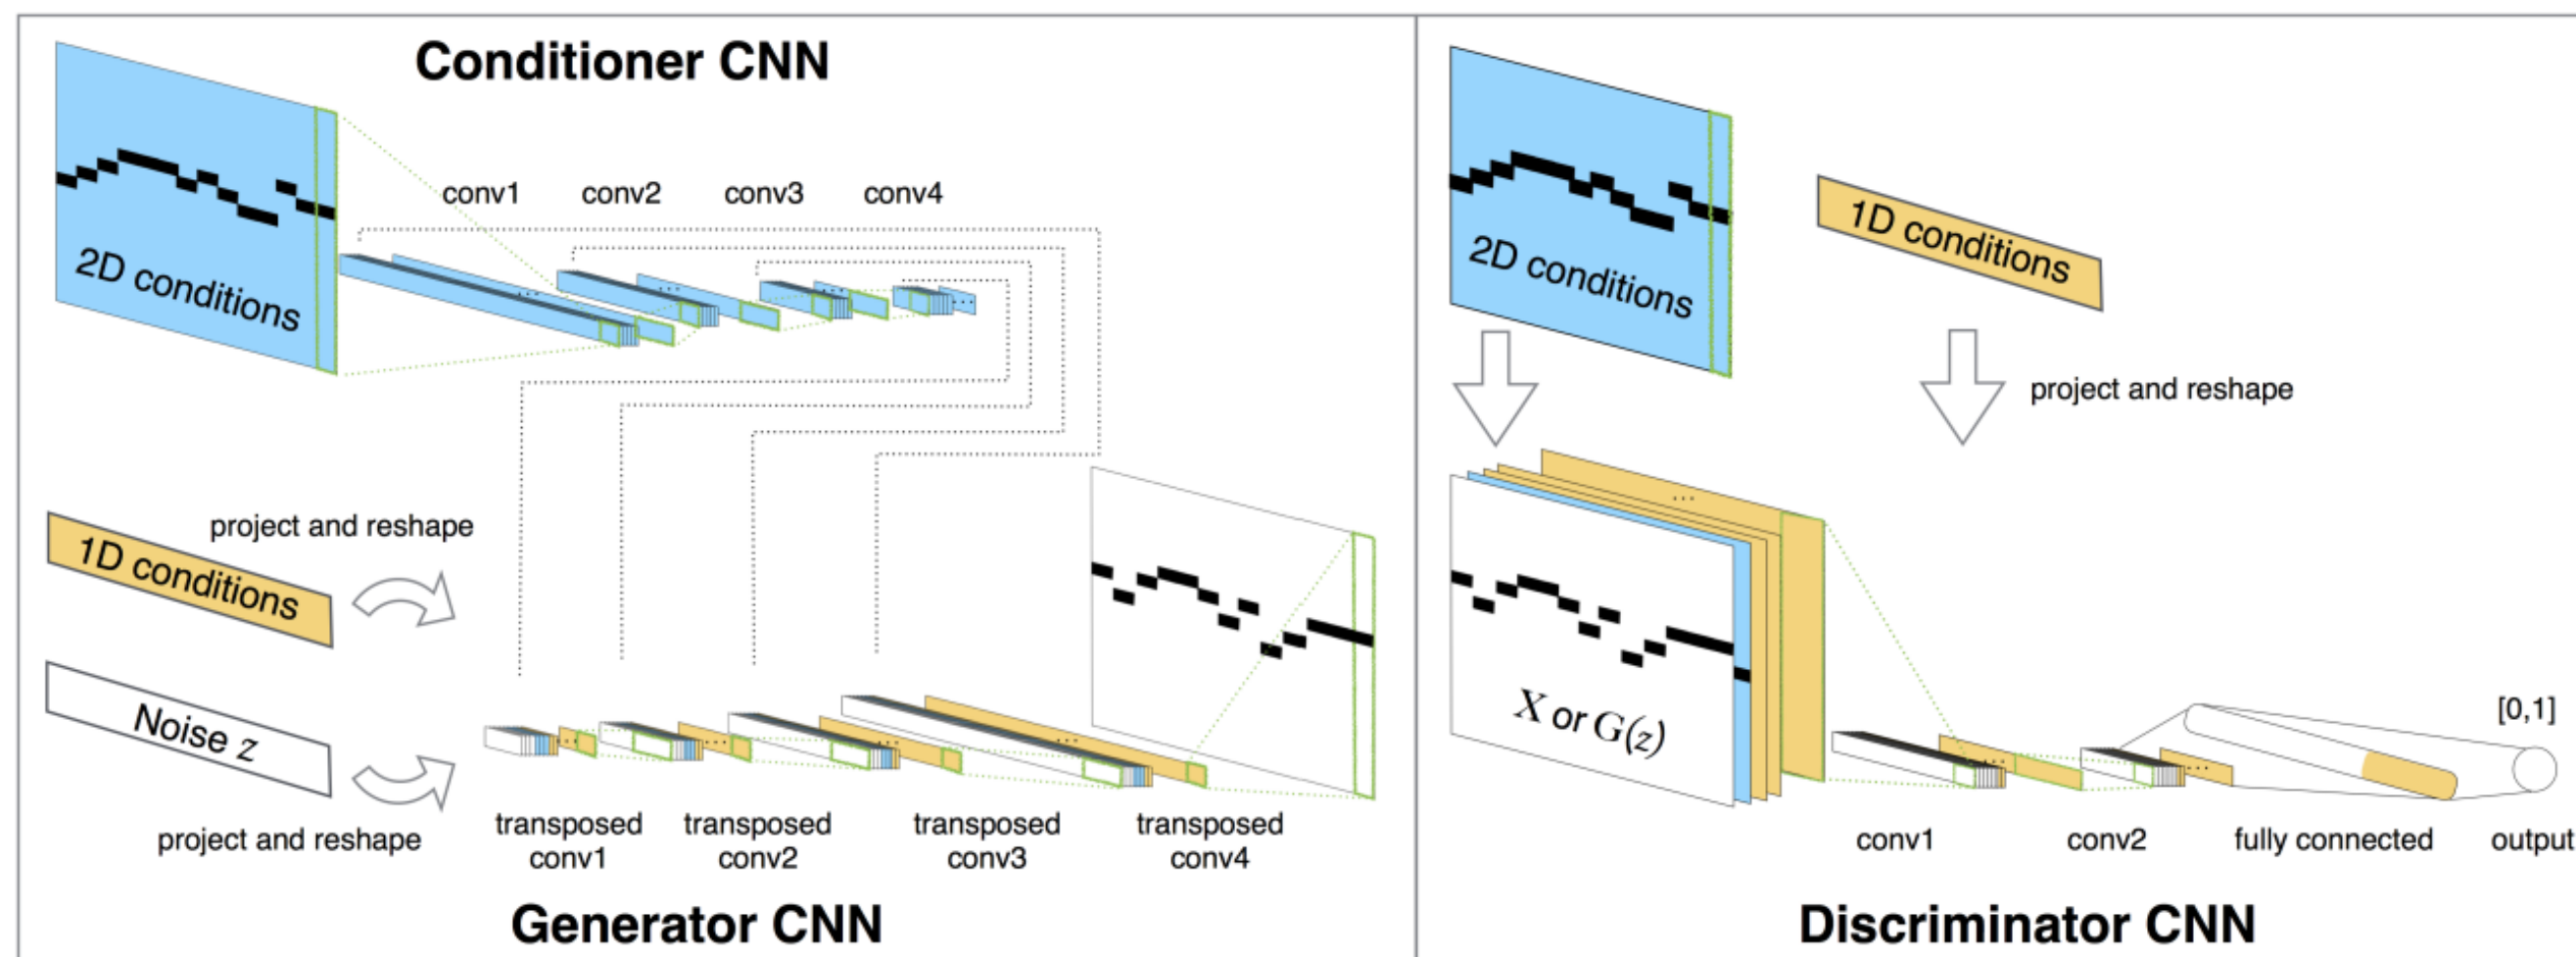
MidiNET:

Convolutional GAN Music Generator

PROBLEM
STATEMENT

PASTWORK

DATA COLLECTION



- h = # of MIDI Notes
- w = # of time steps
- Noise vector z gets fed into the Generator, G .
- Each entry of $G(z)$ gets simultaneously conditioned by 2D conditions (Blue Block)
- Output of $G(z)$ is $h \times w$ matrix.

[HTTPS://ARXIV.ORG/PDF/1703.10847V2.PDF](https://arxiv.org/pdf/1703.10847v2.pdf)

Data Collection



PROBLEM
STATEMENT

PASTWORK

DATACOLLECTION

Data Collection

- **3** Genre Total **18** hours of music
- youtube-dl to scrape music based on genre
- Subset of MAESTRO dataset
- torchaudio to resample and automatically generate melspectrograms from audio, use network to predict next frame of spectrogram
- Allows for **conditional generation** based on genre

Data Pipeline

Scalable **custom MusicDataset** pipeline using Pytorch Dataset class

Divide music into training samples on the scale of **seconds**

Convert it from **wave form** to **MelSpectrogram**

MeINET

RESEARCH

SET UP MODEL

OPTIMIZATION

DIVERGE
TRAINING

ANALYZE
RESULTS

Autoregressive element-by-element music generation

- **RNN** based autoregressive conditional generative model
- **Time-delayed** stack and **Frequency-delayed** stack
- Feature **extraction** networks
- **Gaussian** mixture modelling
- **Multiscale** modelling

$$p(x) = \prod_i \prod_j p(x_{ij} | x_{<ij}; \theta_{ij})$$

$$p(y|x) = \sum_k \pi_k(x) N(y | \mu_k(x), I\sigma_k^2(x)).$$

- **Training objective:** minimizing mixture density networks loss

$$L(w) = \frac{-1}{N} \sum_{n=1}^N \log \left(\sum_k \pi_k(x_n, w) N(y_n | \mu_k(x_n, w), I\sigma_k^2(x_n, w)) \right),$$

MeINET

Implemented in Pytorch

RESEARCH

SET UP MODEL


OPTIMIZATION

DIVERGE
TRAINING


ANALYZE
RESULTS

```
39 class TimeDelayedStack(nn.Module):
40     def __init__(self, dims):
41         super().__init__()
42         self.bi_freq_rnn = nn.GRU(dims, dims, batch_first=True, bidirectional=True)
43         self.time_rnn = nn.GRU(dims, dims, batch_first=True)
44
45     def forward(self, x_time):
46
47         # Batch, Timesteps, Mels, Dims
48         B, T, M, D = x_time.size()
49
50         # Collapse the first two axes
51         time_input = x_time.transpose(1, 2).contiguous().view(-1, T, D)
52         freq_input = x_time.view(-1, M, D)
53
54         # Run through the rnns
55         x_1, _ = self.time_rnn(time_input)
56         x_2_and_3, _ = self.bi_freq_rnn(freq_input)
57
58         # Reshape the first two axes back to original
59         x_1 = x_1.view(B, M, T, D).transpose(1, 2)
60         x_2_and_3 = x_2_and_3.view(B, T, M, 2 * D)
61
62         # And concatenate for output
63         x_time = torch.cat([x_1, x_2_and_3], dim=3)
64         return x_time
```

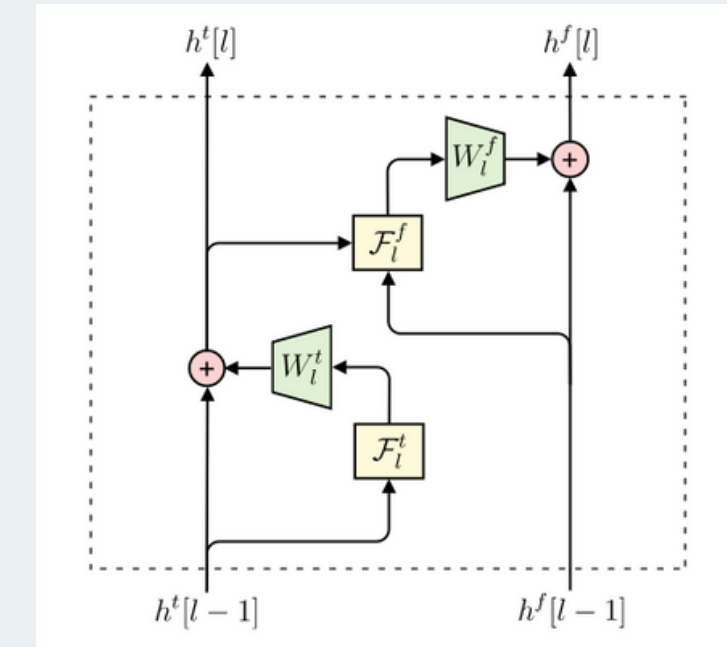
```
20 class FrequencyDelayedStack(nn.Module):
21     def __init__(self, dims):
22         super().__init__()
23         self.rnn = nn.GRU(dims, dims, batch_first=True)
24
25     def forward(self, x_time, x_freq):
26         # sum the inputs
27         x = x_time + x_freq
28
29         # Batch, Timesteps, Mels, Dims
30         B, T, M, D = x.size()
31         # collapse the first two axes
32         x = x.view(-1, M, D)
33
34         # Through the RNN
35         x, _ = self.rnn(x)
36         return x.view(B, T, M, D)
```



(a) Time-delayed stack



(b) Frequency-delayed stack



$$h_{ij}^t[l] = W_l^t F_l^t(h^t[l-1])_{ij} + h_{ij}^t[l-1] .$$

MeINET

RESEARCH


SET UP MODEL

OPTIMIZATION


DIVERGE
TRAINING

ANALYZE
RESULTS


MULTISCALE MODELLING



(a) Tier 1 (32 × 50)



(b) Tiers 1-3 (64 × 100)



(c) Tiers 1-6 (256 × 200)

$$p(x; \psi) = \prod_g p(x^g | x^{<g}; \psi^g).$$

GENERATES SPECTROGRAMS FROM
HIGH-LEVEL STRUCTURE TO FINE-GRAINED DETAILS BY
ITERATING THROUGH THE **MULTIPLE TIERS** THAT A
SPECTROGRAM IS PRE-PARTITIONED INTO

EACH TIER CONTAINING HIGHER RESOLUTION OF
INFORMATION THAN THE PREVIOUS ONE

FEATURE EXTRACTION NETWORKS

```
98 class FeatureExtraction(nn.Module):
99     def __init__(self, num_mels, n_layers):
100         super().__init__()
101         # Input layers
102         #self.freq_fwd = nn.GRU(time_steps, time_steps, batch_first=True)
103         #self.freq_back = nn.GRU(time_steps, time_steps, batch_first=True)
104         self.time_fwd = nn.GRU(num_mels, num_mels, batch_first=True)
105         self.time_back = nn.GRU(num_mels, num_mels, batch_first=True)
106         self.weights = nn.Linear(2,1)
107         #self.num_params()
108
109     def forward(self, spectrogram):
110         # Shift the inputs left for time-delay inputs
111         # spectrogram: (batch_size, time, freq)
112         #print("spec", spectrogram.shape)
113         #N,T,F = spectrogram.size()
114         #freq_input = spectrogram.transpose(1, 2).contiguous()
115         #print("the fatuers extracting from", type(spectrogram), len(spectrogram))
116         time_fwd_feats, _ = self.time_fwd(spectrogram)
117         time_back_feats, _ = self.time_back(spectrogram.flip(1))
118         #freq_fwd_feats, _ = self.freq_fwd(freq_input)
119         #freq_back_feats, _ = self.freq_back(freq_input.flip(2))
120
121         #freq_features = freq_features.transpose(1,2).contiguous().view(-1, T, 2*F)
122         stacked = torch.stack((time_fwd_feats, time_back_feats), dim=-1)#, freq_fwd_feats)
123         #assert not torch.any(torch.isinf(spectrogram))
124         #print(spectrogram.max())
125         #assert not torch.any(torch.isinf(time_fwd_feats))
126         #assert not torch.any(torch.isinf(stacked))
127         #assert not torch.any(torch.isinf(self.weights(stacked)))
128         return self.weights(stacked).squeeze(-1)
129
130     def num_params(self):
131         parameters = filter(lambda p: p.requires_grad, self.parameters())
132         parameters = sum([np.prod(p.size()) for p in parameters]) / 1_000_000
133         print('Trainable Parametersextrac: %.3fM' % parameters)
134
```

MeINET



Initial approach was **far too slow**

Speed up training by a factor of **10**

- **Improved** resampling processing and MelSpectrogram conversion
- Further reduced disk accesses from **$O(N)$ to $O(1)$**
- Used a different **batching method** to expand usable dataset
- Set multiple worker, it is now **multi-threaded** using the dataloader API
- Added buffer checkpointing, **~50 times** increase in model expressive power
- Changed the data pipeline to be optimal on Google Colab

MeINET

RESEARCH

SET UP MODEL

OPTIMIZATION

DIVERGE
TRAINING

ANALYZE
RESULTS

Model 1 (big)

50.8 M Parameters
~ 5hrs / epoch (3.6 Hz)

30 epochs

DataConfig:

batch_sz=6, num_mels=180,
win_sz=3,
stft_hop_sz=800,
stft_win_sz=256*8

TrainConfig:

dims=256, n_layers=[12,6,5,4],
directions=[2,1]

Model 2 (small)

4.04 M Parameters
~ 1hr / epoch (18 Hz)

30 epochs


DataConfig:

batch_sz=6, num_mels=180,
win_sz=8,
stft_hop_sz=800,
stft_win_sz=256*8

TrainConfig:

dims=64, n_layers=[12,6,5,4],
directions=[2,1]

MeINET



Very hard to train, need to build another model.

Comparison results at the end of the presentation.



FIRST OF EVER
WE PROPOSE:

CMelGAN

Inspired by MelGAN and cGAN
Conditional Generative Adversarial Networks
based on Melspectrograms

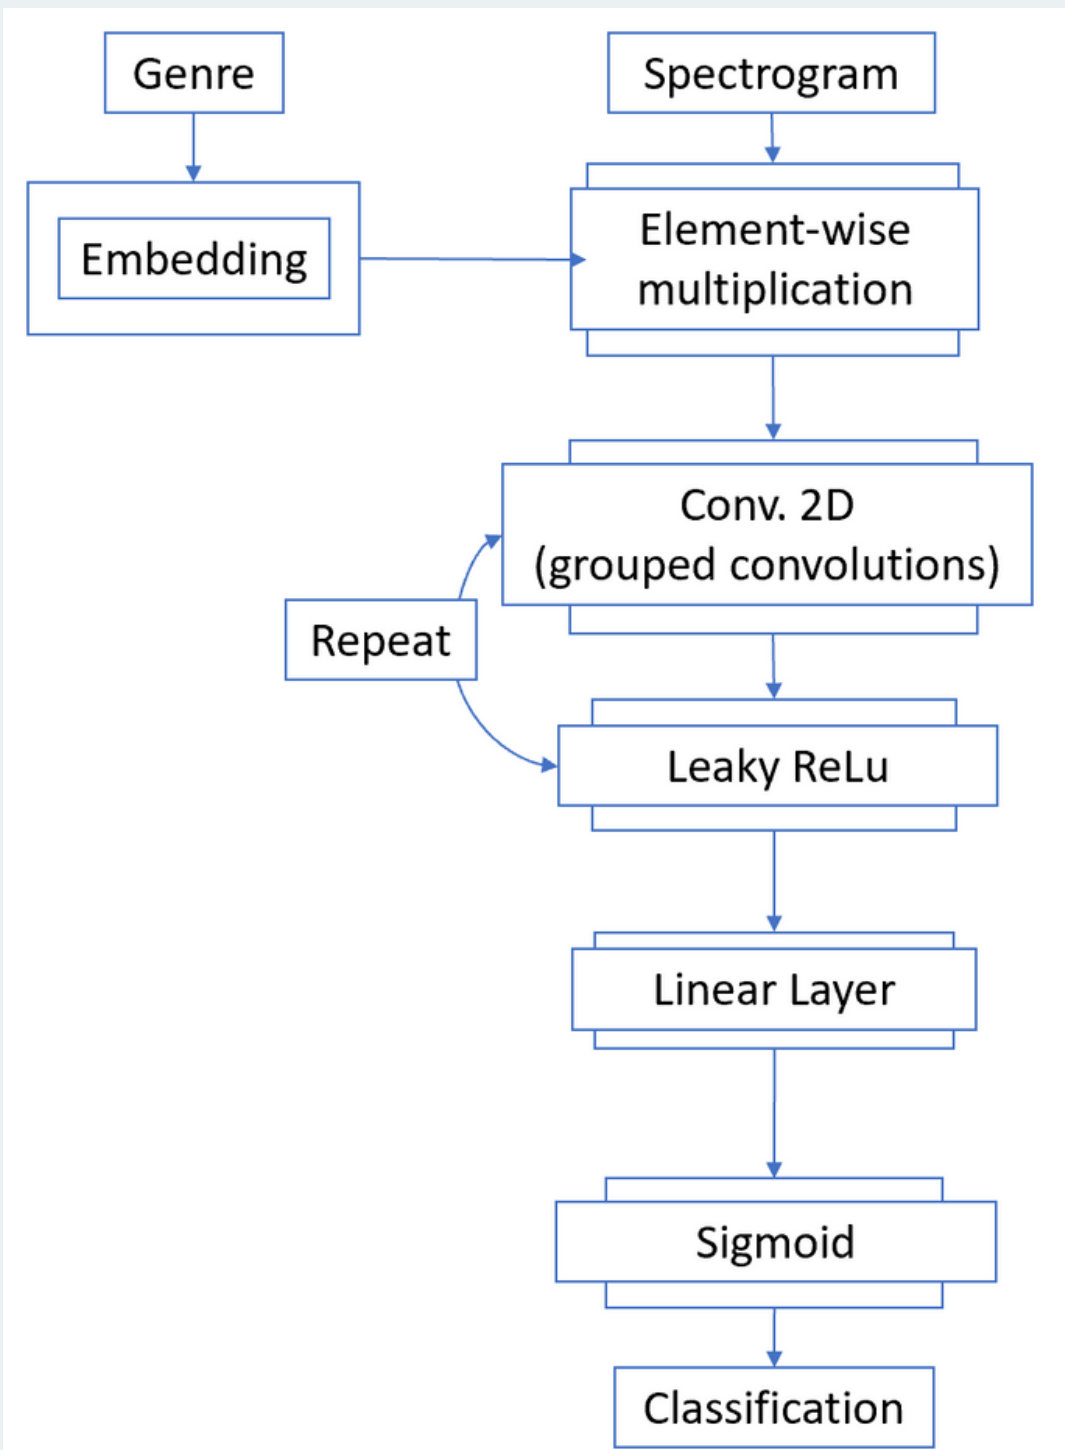
CMelGAN

SET UP MODEL

OPTIMIZATION

ANALYZE RESULTS

Discriminator



Combines ideas from MelGAN and CGAN

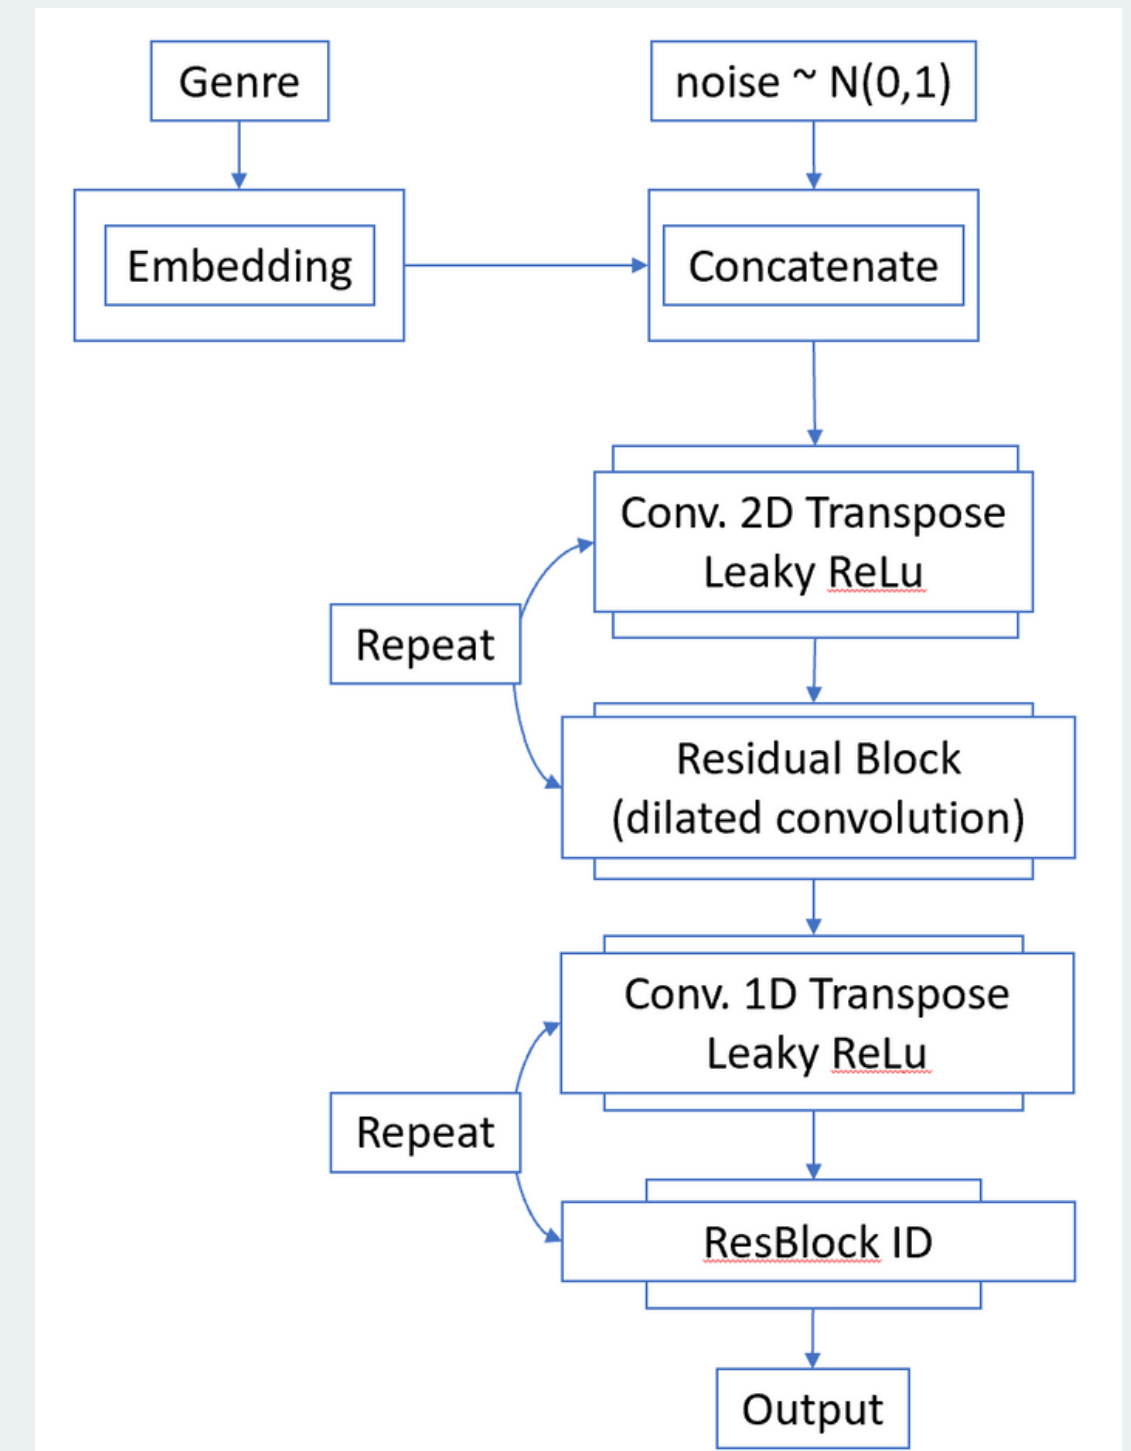
Generator

- **Embedding** layer for genre, elementwise multiplication
- Transposed **2D** convolutions, with weight normalization
- **Dropout** and batchnorm for regularization
- **Residual** blocks with **dilated convolutions**
- Transposed 1D convolutions, with weight normalization
- **Dropout** and 1D batchnorm

Discriminator

- **2D** grouped convolution, linear layer

Generator



CMelGAN



SET UP MODEL

OPTIMIZATION

ANALYZE
RESULTS

Optimization of model

- Run on **tiny fraction** of dataset (~10%) to get quick look at hyperparameter performance
- Most hyperparameter settings were unstable, with either the discriminator either being easily fooled, or the generator not powerful enough to fool the discriminator
- With fractional trick, were able to test far more hyperparameter settings
- Settled on some decent ones that made neither discriminator nor generator dominate
- Add dropout, batchnorm, and weight normalization (bag of tricks)

Preliminary Results

RESEARCH

SET UP MODEL

ANALYZE RESULTS

Model	MOS (/5)	Processing Rate (train)
MelNet (small)	0.0 (random noise)	307 kHz
MelNet (large)	0.0 (random noise)	75 kHz
GriffinLim (MelNet settings)	0.5 (basically random noise)	-
ConditionalMelGAN (medium)	0.0 (random noise)	980 kHz
CMelGAN (large, no 1D Convolution)	0.0 (random noise)	705 kHz
GriffinLim (CMelGAN settings)	0.5 (basically random noise)	-

Conclusion and Next Steps

RESEARCH

SET UP MODEL

ANALYZE
RESULTS

- Melspectrograms are really **hard to work with**
- **Unstructured format** makes it too hard for networks to learn a good representation
- **Griffin-Lim** algorithm not good enough for inversion, need to add a neural vocoder
- MIDI approach seems to be **better option**, at least on **limited hardware**
- Network was **unable** to learn a good representation of music
- Probably underestimated the difficulty of the project

